



Menu



Browse Library



Search in our library...



Training Library / Create and Manage a Kubernetes Cluster from Scratch

# Initializing the Kubernetes Master Node

49m 56s left

## Open Cloud Environment



100%

Setup completed

Average setup time: 3m 29s

## Credentials

Account ID

030925358

Copy

Username

student

Copy

Password

Ca1\_6eFlz

Copy

Region

US West 2

Copy

PEM

PPK

Download

Download



100%

Bridge Connection Completed

## Lab Steps

1

Logging In to Amazon Web Services Console

Skip to content

## Introduction

You will use `kubeadm` to initialize the control-plane node in this lab step. The initialization process will create a certificate authority for secure cluster communication and authentication, and start all the node components (`kubelet`), control-plane components (API server, controller manager, scheduler, `etcd`), and common add-ons (`kube-proxy`, DNS). You will see how easy the initialization process is with `kubeadm`.

The initialization uses sensible default values that adhere to best practices. However, [many command options](#) are available to configure the process, including if you want to provide your own certificate authority or if you want to use an external `etcd` key-value store. One option that you will use is required by the pod network plugin that you will install after the control-plane is initialized. `kubeadm` does not install a default network plugin for you. You will use Calico as the pod network plugin. Calico supports Kubernetes network policies. For network policies to function properly, you must use the `--pod-network-cidr` option to specify a range of IP addresses for the pod network when initializing the control-plane node with `kubeadm`.

There are many network plugins besides Calico. Calico is used primarily because it is used in clusters in [Kubernetes certification exams](#) and it supports network policies. Calico is used internally by AWS, Azure, and GCP for their managed Kubernetes offerings, so you can be certain it is production-ready. However, if all of your environments live in AWS, you may consider the [Amazon VPC network plugin](#).

## Instructions

1. Initialize the control-plane node using the init command:

Copy code

```
1 | sudo kubeadm init --pod-network-cidr=192.168.0.0/16
```

Support

The pod network CIDR block (192.168.0.0/16) is the default used by network CIDR used configuration of

Press `option` + `Q` to open this menu



## Instance Connect

3

Installing kubeadm  
and Its  
Dependencies

4

Initializing the  
Kubernetes Master  
NodeJoining a Worker  
Node to the  
Kubernetes ClusterNeed help? Contact our  
support team

```
[preflight] Running pre-flight checks
[preflight] Pulling images required for setting up a Kubernetes cluster
[preflight] This might take a minute or two, depending on the speed of your internet connection
[preflight] You can also perform this action in beforehand using 'kubeadm config images pull'
[certs] Using certificateDir folder "/etc/kubernetes/pki"
[certs] Generating "ca" certificate and key
[certs] Generating "apiserver" certificate and key
[certs] apiserver serving cert is signed for DNS names [ip-10-0-0-90.kubernetes.kubernetes.default.kubernetes.default.svc.kubernetes.default.cluster.local] and IPs [10.96.0.1 10.0.0.90]
[certs] Generating "apiserver-kubelet-client" certificate and key
[certs] Generating "front-proxy-ca" certificate and key
[certs] Generating "front-proxy-client" certificate and key
[certs] Generating "etcd/ca" certificate and key
[certs] Generating "etcd/server" certificate and key
[certs] etcd/server serving cert is signed for DNS names [ip-10-0-0-90.localhost] and IPs [10.0.0.90 127.0.0.1 ::1]
[certs] Generating "etcd/peer" certificate and key
[certs] etcd/peer serving cert is signed for DNS names [ip-10-0-0-90.localhost] and IPs [10.0.0.90 127.0.0.1 ::1]
[certs] Generating "etcd/healthcheck-client" certificate and key
[certs] Generating "apiserver-etcd-client" certificate and key
[certs] Generating "sa" key and public key
[kubeconfig] Using kubeconfig folder "/etc/kubernetes"
[kubeconfig] Writing "admin.conf" kubeconfig file
[kubeconfig] Writing "kubelet.conf" kubeconfig file
[kubeconfig] Writing "controller-manager.conf" kubeconfig file
[kubeconfig] Writing "scheduler.conf" kubeconfig file
```

Read through the output to understand what is happening. At the end of the output, useful commands for configuring `kubectl` and joining worker nodes to the cluster are given:

```
To start using your cluster, you need to run the following as a regular user:

mkdir -p $HOME/.kube
sudo cp -i /etc/kubernetes/admin.conf $HOME/.kube/config
sudo chown $(id -u):$(id -g) $HOME/.kube/config

Alternatively, if you are the root user, you can run:

export KUBECONFIG=/etc/kubernetes/admin.conf

You should now deploy a pod network to the cluster.
Run "kubectl apply -f [podnetwork].yaml" with one of the options listed at:
https://kubernetes.io/docs/concepts/cluster-administration/addons/

Then you can join any number of worker nodes by running the following on each as root:

kubeadm join 10.0.0.92:6443 --token c5kqhz.2h8efb7qmxu3kzx9 \
--discovery-token-ca-cert-hash sha256:1cd9c60882f9c9d7630453ae884606d76311e7fc
```

2. Copy the `kubeadm join` command at the end of the output and store it somewhere you can access later.

It is simply convenient to reuse the given command, although you can regenerate it and create new tokens using the `kubeadm token` command. The join tokens expire after 24 hours by default.

3. Initialize your user's default `kubectl` configuration using the `admin` `kubeconfig` file generated by `kubeadm`:

Copy code

```
1 mkdir -p $HOME/.kube
2 sudo cp -i /etc/kubernetes/admin.conf $HOME/.kube/config
3 sudo chown $(id -u):$(id -g) $HOME/.kube/config
```

4. Confirm you can use `kubectl` to get the cluster component statuses:

Skip to content

Press + to open this menu

Copy code



```
etcd-0           Healthy  {"health":"true","reason":""}
controller-manager Healthy  ok
scheduler        Healthy  ok
```

The output confirms that the **scheduler**, **controller-manager**, and **etcd** are all **Healthy**. The Kubernetes API server is also operational, or `kubectl` would have returned an error attempting to connect to the API server.

Enter `kubeadm token --help` if you would like to know more about `kubeadm` tokens.

5. Get the nodes in the cluster:

[Copy code](#)

```
1 | kubectl get nodes
```

NAME	STATUS	ROLES	AGE
ip-10-0-0-92	NotReady	control-plane	7m24s

The control-plane node is reporting a **STATUS** of **NotReady**. Notice `kubeadm` gives the node a **NAME** based on its IP address. The `--node-name` option can be used to override the default behavior.

6. Describe the node to probe deeper into its **NotReady** status:

[Copy code](#)

```
1 | kubectl describe nodes
```

In the **Conditions** section of the output, observe the **Ready** condition is **False**, and read the **Message**:

```
NetworkReady=false reason:NetworkPluginNotReady message:Network plugin returns error: cni plugin not initialized
```

The kubelet is not ready because the network plugin is not ready. The **cni config uninitialized** refers to the [container network interface](#) (CNI) and is a related problem. Network plugins implement the CNI interface. You will resolve the issue by initializing the Calico network plugin.

7. Enter the following commands to create the Calico network plugin for pod networking:

[Copy code](#)



```
customresourcedefinition.apiextensions.k8s.io/globalnetworkpolicies.crd.projectcalico.org created
customresourcedefinition.apiextensions.k8s.io/globalnetworksets.crd.projectcalico.org created
customresourcedefinition.apiextensions.k8s.io/hostendpoints.crd.projectcalico.org created
customresourcedefinition.apiextensions.k8s.io/ipamblocks.crd.projectcalico.org created
customresourcedefinition.apiextensions.k8s.io/ipamconfigs.crd.projectcalico.org created
customresourcedefinition.apiextensions.k8s.io/ipamhandles.crd.projectcalico.org created
customresourcedefinition.apiextensions.k8s.io/ippools.crd.projectcalico.org created
customresourcedefinition.apiextensions.k8s.io/kubecontrollersconfigurations.crd.projectcalico.org created
customresourcedefinition.apiextensions.k8s.io/networkpolicies.crd.projectcalico.org created
customresourcedefinition.apiextensions.k8s.io/networksets.crd.projectcalico.org created
clusterrole.rbac.authorization.k8s.io/calico-kube-controllers created
clusterrolebinding.rbac.authorization.k8s.io/calico-kube-controllers created
clusterrole.rbac.authorization.k8s.io/calico-node created
clusterrolebinding.rbac.authorization.k8s.io/calico-node created
daemonset.apps/calico-node created
serviceaccount/calico-node created
deployment.apps/calico-kube-controllers created
serviceaccount/calico-kube-controllers created
poddisruptionbudget.policy/calico-kube-controllers created
```

A variety of resources are created to support pod networking. A daemonset is used to run a Calico-node pod on each node in the cluster. The resources include several custom resources (customresourcedefinition) that extend the Kubernetes API, for example, to support network policies (networkpolicies.crd.projectcalico.org). Many network plugins have a similar installation procedure.

8. Watch the status of the nodes in the cluster:

 Copy code

```
1 | watch kubectl get nodes
```

NAME	STATUS	ROLES	AGE
ip-10-0-0-11	Ready	control-plane	2m41s

With the network plugin initialized, the control-plane node is now **Ready**.

*Note:* It may take a minute to reach the **Ready** state.

Press **ctrl+c** to stop watching the nodes.

## Summary

In this lab step, you used `kubeadm` to initialize a control-plane node. You also initialized a pod network plugin named Calico to fully bring up the control-plane node. The EC2 instance type used in the lab does not have enough CPU capacity to satisfy the CPU resource requests of all of the pods on one instance. The `kube-dns` pod is currently unschedulable due to a lack of CPU resources, although the cluster can operate without it. Once you join a worker node to the cluster in the next lab step, there will be enough CPU capacity for the `kube-dns` pod.

It is worth mentioning that with a single control-plane, there is a single point

Skip to content

Press  +  to open this menu

ut you should  
he procedure is  
balancer to distribute

## VALIDATION CHECKS

### 1 Checks

[Start check](#)

#### Created Calico Network Plugin

Check for calico pods to confirm creation.

Kubernetes

Did you like  
this step?

[× End Lab](#)[< Back](#)[Start check](#)

## ABOUT US

[About Cloud Academy](#)[About QA](#)[About Circus Street](#)

## COMMUNITY

[Join Discord Channel](#)

## HELP

[Help Center](#)